D · What the actual agent trace reveals

One recorded Gemini run · structural errors, publication latency and model validation are separate measures

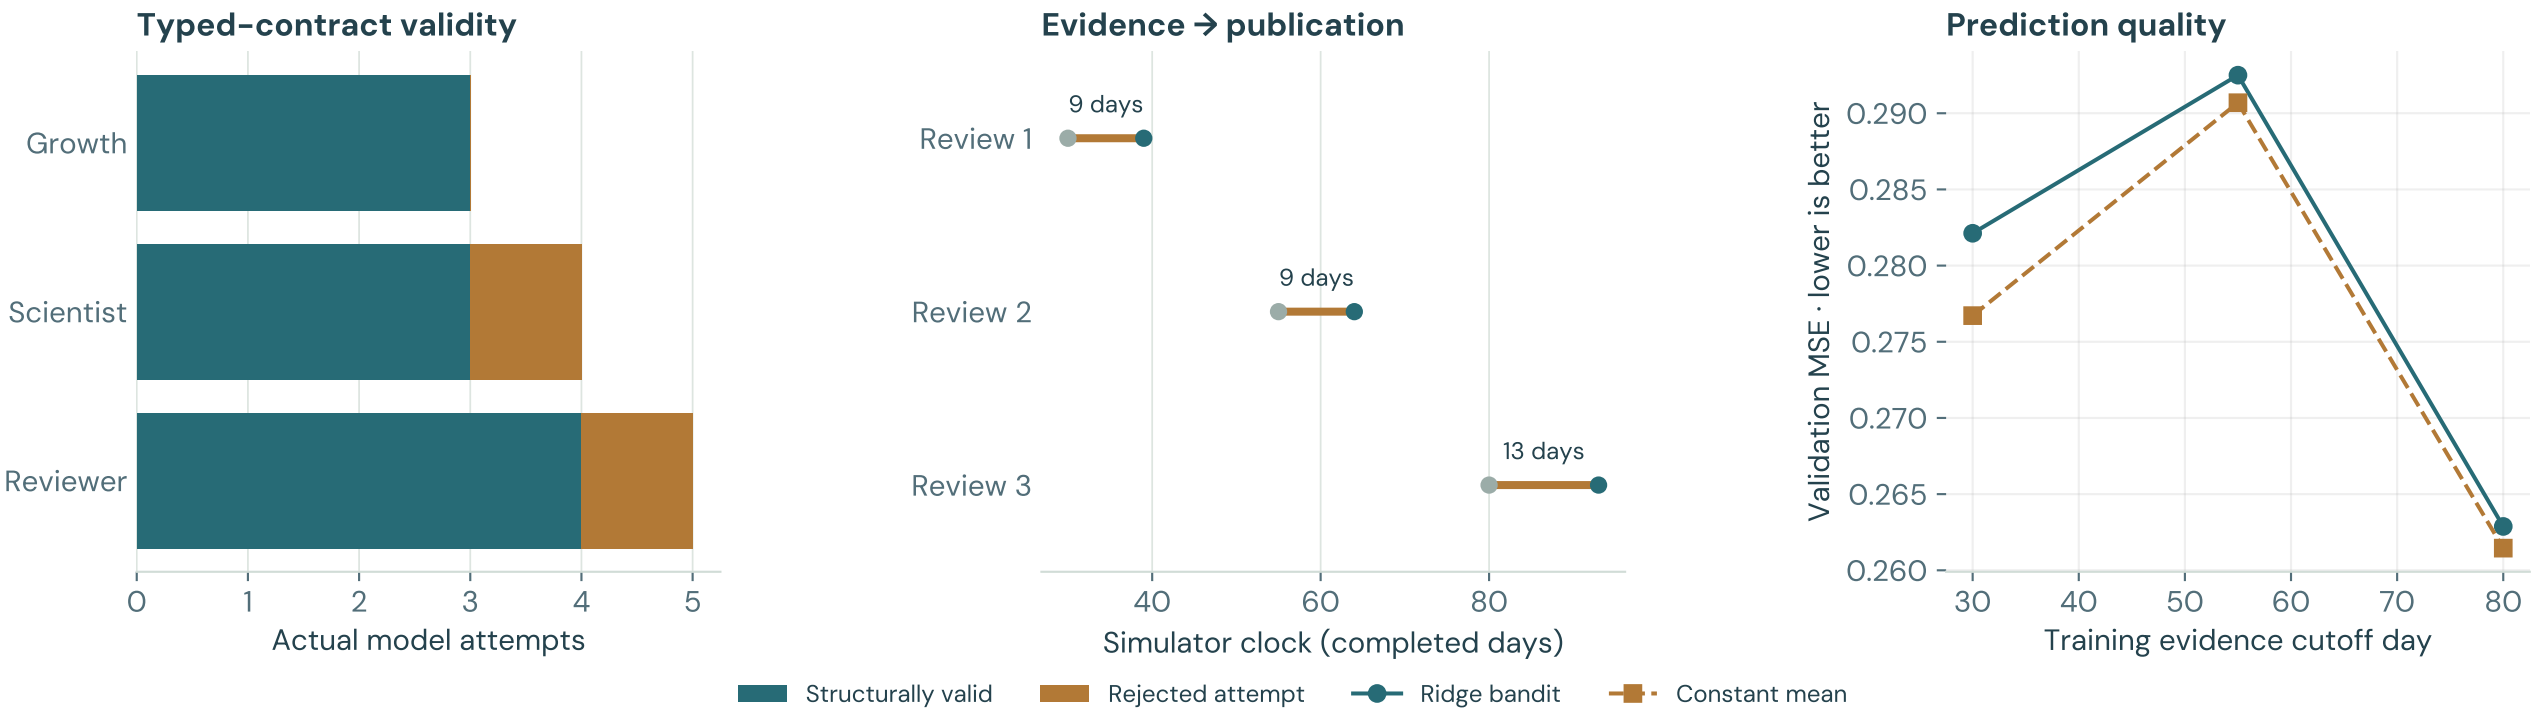

12 model attempts · 2 structurally rejected · $0.008447 recorded token estimate
Four non-exhaustive manual semantic findings remain visible. Contract acceptance is not factual accuracy; MSE is not campaign lift.

Source: GARDEN-LAB-df43cd08499c0c0a · no new model calls in the evaluation suite